

Homework 1

1.2

Environment	Return Mean	Return Std
Ant	4765.671875	101.028511
Humanoid	250.591323	35.290046

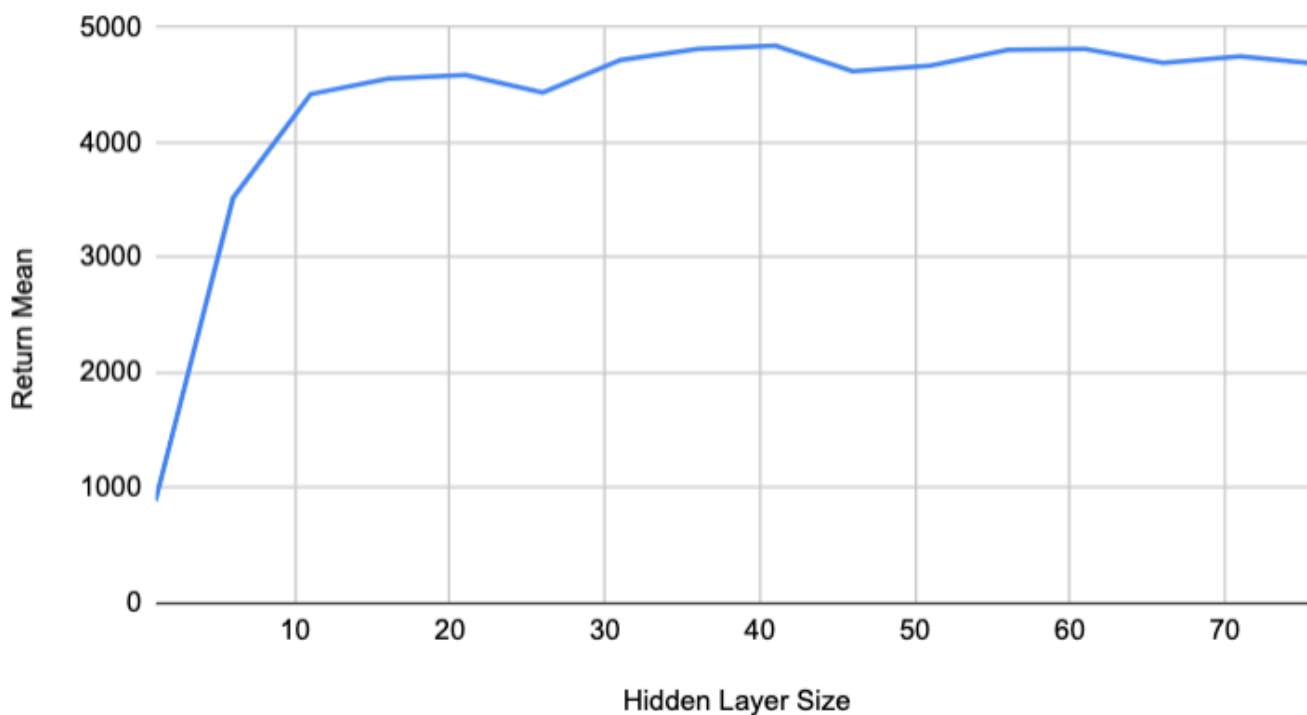
In comparison, the expert achieved an average return of 4713.653320 for Ant and 10344.517578 for Humanoid, with standard deviations of 12 and 21, respectively. Therefore behavioral cloning for Ant produced almost 100% of the average return of the expert, while behavioral cloning only produced about 1% of the average return of the expert.

These results were collected with `ep_len = 1000` and `eval_batch_size = 5000`. All of the other hyperparameters (hidden layer size, number of layers, learning rate, etc.) were the defaults specified in `run_hw1.py`.

1.3

These results were collected for the Ant environment.


BC Agent Performance vs. Hidden Layer Size




Keeping all other default parameters the same, I decided to perform a sweep across the hyperparameter corresponding to hidden layer size. I chose this parameter because I was curious if the default, 64 was too small, or if by increasing the hidden layer size (thus increasing the total number of parameters of the model) we could increase behavioral cloning performance. I was surprised to see that the hidden layer size need only be greater than 10 in order to achieve competitive results with larger sizes. Increasing the size without increasing the number of training steps did not increase performance, though it also did not hurt performance.

2.2

DAgger Return Mean vs Number of Iterations for Ant Environment



DAgger Return Mean vs Number of Iterations for Humanoid Environment



For both of these experiments, I used `ep_len = 1000` and `eval_batch_size = 5000`. All of the other default parameters were the same.